

Media Monitoring

Tracking press and social-media coverage of the Advisory Board and its reports

What the module does today · how the pipelines work · current limitations and next steps

Status overview — 2 August 2026

Prepared from the MethodHub codebase (functional review of the media-monitoring module)

What the module is

A dashboard under /beta/media-monitoring that tracks where and how the Advisory Board, its people and its reports appear in European press and on social media. It answers three questions:

Who covers us?

Which outlets, countries and languages pick up ESABCC work

What resonates?

Which reports and keywords drive the coverage

How far does it go?

Indicative audience reach and tier of the covering outlets

129

search keywords

12

EU languages

120

curated outlets

13

ESABCC reports
tracked

2

channels: press +
social


39

blocked spam domains

Two data channels feed one dashboard:

- **Press** — articles found via Google News RSS search across 129 multilingual keywords, filtered to a curated registry of 120 European and international outlets.
- **Social** — LinkedIn posts captured with a purpose-built browser extension or pasted manually; optional RSS sources can be added per author.

How press coverage is collected



Key mechanics

- **Smart querying** — short keywords and acronyms are exact-matched in quotes; longer phrases match broadly. Up to 5 keywords share one OR query to stay within Google rate limits (3 concurrent requests, throttled waves).
- **Deduplication & attribution** — the real outlet URL is extracted from the Google redirect, tracking parameters stripped, and articles merged across keywords; every article records which keywords and which ESABCC reports it matches.
- **Run audit** — every fetch writes a `media_fetch_runs` row (trigger, articles found/new, errors); overlapping runs are rejected and repeat runs are rate-limited (409/429 guards).
- **Trigger** — a "Refresh feeds" button on the dashboard, or a secret-protected GET endpoint intended for an external scheduler.

Quality controls on the press feed

What gets IN

- **Curated outlet registry** — 120 outlets across ~30 countries — pan-EU trade press (POLITICO, EURACTIV, Carbon Brief...), national dailies in 12 languages, and global wires — each with country, tier, audience estimate and reach score.
- **Every article enriched** — outlet, country, language, tier, estimated audience, matched keywords and matched ESABCC reports.

What stays OUT

- **Blocklist** — 39 domains: press-release wires, SEO aggregators, finance-content farms and known climate-sceptic outlets are dropped at ingest.
- **Unknown outlets** — by default only registry outlets are kept, so the long tail of blogs Google News surfaces never reaches the database.
- **Duplicates** — the same story syndicated at several URLs is merged by canonical URL.

Reach & tier model

- **Tiers** — global · national · trade · regional · blog — used for the tier-mix chart and the tier filter in the article list.
- **Reach** — each outlet carries an estimated monthly audience (public Similarweb/Press Gazette figures) and a 0–10 reach score. Note: this is outlet-level audience, used as an indicative weight — not measured readership of the individual article (see limitations).

The social channel (LinkedIn-first)

Browser extension

A Chrome/Edge (Manifest V3) extension for the comms team: right-click or Alt+Shift+C on any LinkedIn post captures URL, text, author and timestamp from the signed-in session and sends it to the ingest API. No scraping infrastructure needed.

Manual entry

A paste form on the Social tab: drop in a post URL and its text; the server extracts the author and stores it like any captured post.

RSS sources (optional)

Per-author feeds (e.g. via bridge services) can be registered and pulled on demand; sources can be flagged as board members to enable the board-member filter.

What happens at ingest

- **Report matching** — post text is matched against each report's terms, so posts appear on the per-report coverage pages next to the press articles.
- **Deduplication** — posts are upserted on the post URL — capturing the same post twice updates rather than duplicates.
- **Access control** — the ingest endpoint requires the shared MEDIA_MONITORING_SECRET (fail-closed); the extension stores the key in its options page.
- **Engagement metrics** — like/share counts are deliberately not scraped; reach is not tracked for social posts today.

The dashboard — five tabs

Overview

KPI tiles (articles, outlets, countries, total reach), coverage timeline, Europe map of covering outlets, top outlets/countries, keyword frequency, tier and language mix. Adjustable 7–180 day window. Press channel only today.

Reports

One card per ESABCC report: press articles, social posts and indicative reach in the window, with a detail page per report (timeline, top outlets, full coverage list).

Articles

Filterable, sortable article list — search, country, outlet, tier, known-outlets-only, per-report filter; links out to the original pieces.

Social

Captured LinkedIn posts with author/date/report filters, the manual capture form, and RSS source management.

Keywords

The 129 search terms grouped by category with per-keyword article counts, activate/deactivate toggles, add/remove — plus the admin key for protected actions.

Coverage per ESABCC report

Every article and post is matched at ingest against the report catalogue (13 reports, each with 2–4 match terms). The Reports tab turns this into per-report coverage pages — the closest thing the Board has to an impact view per publication.

Towards EU climate neutrality (2024)

progress report

Scientific advice for EU 2040 target

advisory

ETS review reports

advisory

CAP / agri-food reports

advisory

Adaptation reports (2025/2026)

advisory

...and 8 more

full catalogue in `src/data/esabcc-reports.ts`

Each report detail page shows

- A press/social/combined coverage timeline since publication
- Top covering outlets with tier and audience estimates
- The complete article and post list, deep-linked to the sources

Operations & security (hardened Aug 2026)

Operating model

- **Refresh** — on-demand from the dashboard; cron-ready GET endpoints exist for both channels (secret-protected, force-flag for schedulers).
- **Run guards** — concurrent runs are rejected (409) and rapid re-runs rate-limited (429) so the pipeline cannot be hammered.
- **Audit trail** — every run logged in `media_fetch_runs` with trigger, counts and errors; shown as "last refreshed" in the UI.
- **Cost** — no paid APIs — Google News RSS and user-initiated capture only.

Security model

- **Shared secret** — `MEDIA_MONITORING_SECRET` protects cron triggers and all write endpoints; the dashboard holds it as a locally-stored admin key.
- **Fail-closed ingest** — the extension/manual capture endpoint refuses writes until the secret is configured (503) and caps payloads at 200 KB.
- **SSRF protection** — registered feed URLs are validated at write and fetch time against private networks and metadata endpoints.
- **Injection hardening** — keyword language/country values are strictly validated and URL-encoded before reaching Google News.

Honest status — current limitations

No scheduler is wired up

nothing in the repo calls the cron endpoints; coverage only refreshes when someone clicks "Refresh feeds". The daily automated news commits belong to a different module.

"Reach" overstates

reach figures are outlet monthly audience applied to every article, then summed — indicative of outlet weight, not actual readership of the coverage.

Keyword stats can inflate

articles found by batched OR queries are credited to all 5 batched keywords when no literal match is found in the RSS snippet.

Old articles are never re-matched

report matching runs at ingest only; articles collected before a report was added to the catalogue never gain that attribution.

Social starts empty

no RSS sources are seeded, and without MEDIA_MONITORING_SECRET set in production the extension/manual capture path refuses writes — the social channel then collects nothing.

Overview is press-only

social activity is invisible on the Overview tab; the Social tab also only displays LinkedIn even though other platforms can be stored.

The allowlist is the bottleneck

only the 120 registry outlets are kept, so genuine coverage in smaller/regional outlets is silently dropped — the main reason volumes look thin.

Navigation label out of date

the hub card still says "GDELT" although the pipeline is Google News RSS.

Recommended next steps

1

Schedule the pipeline

Add a GitHub Action (or Vercel cron) calling both fetch endpoints daily with the secret — turns the module from on-demand into true monitoring.

2

Configure the secret & extension

Set MEDIA_MONITORING_SECRET in production, distribute it to the comms team via the extension options — unlocks social capture.

3

Fix the reach story

Rename to "audience-weighted coverage", or replace with an article-level estimate; stop summing monthly uniques.

4

Bring social into Overview

Add posts to the analytics endpoint and KPI tiles so one screen shows both channels.

5

Add a re-match job

A small script to recompute report attribution over the archive whenever the report catalogue changes.

6

Review the outlet registry

Decide between strict curation and broader capture (e.g. keep unknown outlets flagged for triage instead of dropping them).