

EU Policy Targets Register (M-36)

What we learned from the reviewer training

Lessons extracted from the reviewer prompts and marked-up Excel workbooks supplied across four review rounds, July–August 2026

This document distils the corrections the ESABCC reviewers supplied — through written prompts and mark-ups of the exported workbooks *eupolicytargets_v2_28072026.xlsx* and *eupolicytargets_v3_29072026.xlsx* (reviewers James, Kamila, Mar, Sebastian) — into the general rules now built into the register's extraction and classification pipeline. It is written so a new reviewer, or a future extraction round, can apply the same standards without re-deriving them from the raw correction files.

Round	Input	Register moved
1 · Fact-check (Jul 2026)	Internal audit of all 662 rows, 38 acts; one reviewer agent per act, adversarially re-verified	662 → 653 rows; 312 rows corrected
2 · Expansion + lens (Jul 2026)	24 acts added (resilience, health, cohesion, TEN acquis); relevance lens introduced	1 018 rows / 62 acts; 673 relevant
3 · Human review v2 (Jul 2026)	Reviewer mark-up: 21 rows flagged “not correct” with written reasons, generalised register-wide; then a full-register audit	1 018 → 918 rows; 645 in default lens
4 · Human review v3 (Aug 2026)	Revision sheet: 105 “not a target” marks, 41 comments for AI, 134 first/second-order labels; consolidated-text refresh	918 → 819 targets / 61 acts; 544 in default lens

Sources: docs-internal/policy-targets-factcheck-2026-07.md, policy-targets-human-review-2026-07.md, policy-targets-human-review-2026-08.md; scripts/policy-targets-overrides.json; scripts/policy-targets-input/_corrections-*.json.

1 · The central lesson: what a target is

The single biggest source of error across all rounds was rows that quote real legal text but are not **targets**. The reviewers' corrections converged on one working definition, quoted from the v3 review request:

“Anything that requires just one intervention (like banning something) rather than a progression over time towards a goal is not a target. Targets should be either timebound or at least imply a progression of efforts which can be measured.”

Everything else in this document is that principle applied to specific legal forms. A useful mental test taken from the training: ask **what would progress toward this look like next year?** If the answer is “nothing — it is either complied with or not”, it is an obligation, criterion or ban, not a target.

1.1 · The fifteen “not a target” classes

Each reviewer mark-up was read as an instance of a general class, and the class was applied to every row in the register — the reviewers explicitly said their mark-up was not exhaustive and asked for generalisation. Rules NT-1–NT-6 come from the July round, NT-7–NT-15 from the August round.

Rule	Not a target when the row is...	Reviewer's archetype
NT-1	A content requirement for a plan, strategy or report	Social Climate Fund Art. 4(4): “not a target but a description of what the plans ... should include”
NT-2	An eligibility or funding-scope clause	SCF Art. 7(2)(d): “part of a list of measures eligible to be covered by the fund”
NT-3	An assessment, award or project-selection criterion	TEN-E Art. 4(3): “not a target but a characteristic of a project”
NT-4	A definition or definitional chapeau	EU Taxonomy Art. 9: “a clarification of definitions”
NT-5	A pointer to a target set by another instrument	TEN-E Art. 4(2)(e)(iii): characteristics “in line with targets from other policies ... should not be captured”
NT-6	Context, an illustrative statistic, or a heading	Zero Pollution Action Plan: “just a summary of other commitments as context”
NT-7	A pure conduct obligation	Methane Reg. Art. 13: “an obligation ... not a target to achieve something in the future”
NT-8	A threshold-triggered obligation	EED enterprise energy-audit duties
NT-9	A derogation or reduced requirement	AFIR Arts. 3(8), 4(5): 50 % power-output reductions on low-traffic roads
NT-10	A bare fragment with no actor or date	PPWR recycled-content percentages quoted without their chapeau
NT-11	A qualifying condition on a headline target	NRL rows that “qualify how the headline target should be achieved”
NT-12	A one-off ban or market restriction	PPWR Art. 25(1) packaging-format ban from 2030
NT-13	A methodology or compliance calculation	RED: “part of the methodology for calculating compliance”
NT-14	A Commission assessment or evaluation duty	RED: “an obligation for the Commission to do an evaluation”
NT-15	A process obligation to set objectives later	PPWR Art. 49: Member States to set collection objectives by 2029

Two mechanical companions: **DUP** (the same sentence already registered under a more precise provision of the same act) and **FIX** (a truncated quote superseded by a complete re-extraction, §2 below).

1.2 · Where the rules must NOT be applied — the limits

The rules cut close to real targets, so the training also taught where to stop. These limits were applied consistently and matter as much as the rules:

- **A plan-content item that carries its own date or quantity stays.** EPBD Art. 3(2) (national 2030/2040/2050 renovation-rate targets inside a roadmap requirement) and the MSFD's dated milestones are content requirements in form, but each fixes an outcome and a date. The flagged Social Climate Fund items fixed neither.
- **An obligation to *establish* something stays; a description of what it must *contain* goes.** Floods Directive Art. 7(1) (establish flood risk management plans) is kept; Art. 7(3) (what those plans shall address) is removed.
- **Review and reporting clauses with their own deadline are kept** (CO₂ cars Art. 15, Euro 7 Art. 18, LULUCF Art. 4(4)) — labelled as commitments, not targets.
- **Never drop a row merely for being non-climate.** Relevance is handled by the lens (§5), not by deletion.

2 · Quote extraction: verbatim, complete, in context

The register's standing rule — every quote must be a verbatim, contiguous substring of the act's enacting terms — held in every audit (all 662 first-round quotes were exact; no fabrication was ever found). The failures were subtler: quotes that were *accurate but incomplete*, which is enough to change their meaning. Three truncation modes were identified, all fixed by re-extracting from the EUR-Lex source rather than by editing text:

Mode	Failure	Example fixed
FIX-COLON	A chapeau cut off at its colon, dropping the enumerated points that carry the actual numbers	FuelEU Art. 4(2) now carries the –2 %...–80 % schedule; F-gas Annex VII quota table; CO ₂ -cars Art. 4(1); Taxonomy Arts. 10–11
FIX-ANAPHORA	A sentence opening on a pronoun with no antecedent	"It will aim to ensuring that, by 2030, 75 % of soils are healthy" now includes the Soil Health Mission sentence that governs "It"
FIX-TRUNCATION	A quote sliced mid-sentence — subject before the slice, or the sentence continuing past it	"...will make it mandatory to consider the future climate hazards" (subject: the 2026 Eurocodes update); TEN-T Arts. 41(1)(c), 46(1)

The sharpest lesson on context: **a chapeau can invert a quote's meaning.** ERDF Art. 7(1)(b) reads "investment to achieve the reduction of greenhouse gas emissions..." — but it sits under "*Exclusion from the scope ... The ERDF and the Cohesion Fund shall not support:*". As a row it asserted the opposite of the law. Any quote sitting under "shall include:", "shall contain:", "shall not support:", "means:" or an award-criteria chapeau must be read with that chapeau before it is accepted.

Completeness cuts the other way too:

- **Sibling sub-points.** AFIR Art. 9(1)(a) (90 % shore-side electricity for container ships) was registered while its identical-duty siblings (b) and (c) were not. A "sibling registered, this one not" sweep over quantified sub-points is now standard (added rows in NRL Arts. 4, 5, 14; EED Art. 8(1)(b); 8th EAP Art. 3(s)).
- **Grouped extraction.** Where one provision sets a staged trajectory across sub-points (seven NRL provisions), the reviewers preferred a single row quoting the full provision over one row per stage — the quote ceiling was raised to 2 000 characters to admit them.
- **Commitment lists in communications.** The EU Adaptation Strategy's operative content is fourteen "The Commission will:" bullet lists; the register initially held 8 bullets and missed 40. Soft-law operative bullets are rows.

3 · Scope calls: annexes, duplicates, restatements

- **Annexes that list TARGETS are in; annexes that list ACTIONS are out.** Water Resilience Annex II and Zero Pollution Annex 2 collate targets — captured in full, each row cross-referenced to the act it restates. The Renovation Wave annex (“key Commission actions and indicative timelines”), the Green Deal roadmap annex and Water Resilience Annex I are indexes of actions already described in the body — 20 proposed additions were rejected on this ground.
- **Within one act, duplicates are removed** — including the hard case where one row is a sentence *inside* another rather than a near-identical twin (five such in EPBD Art. 9 alone). Every DUP claim is re-checked mechanically by text containment before removal.
- **Across acts, duplicates are flagged, not removed.** The *duplicate_of* column records the closest target text in another policy (≥ 0.45 token similarity = “duplicate wording”, ≥ 0.25 = “similar target”; a named act in the text wins as “cited in the text”). Thirteen Water Resilience rows restating other instruments’ targets are kept with the cross-reference.
- **A proposed addition that already exists as a peripheral row is a relevance question, not a missing row** — flip the lens flag instead of adding a duplicate.

4 · Classification: substance over keywords

4.1 · Obligation and type

- **Soft law is always voluntary.** Communications and strategies never carry a mandatory obligation, whatever their wording.
- **“Shall endeavour / aim / strive” is voluntary**, even inside a regulation (PPWR 2040 endeavour sub-targets, CAP Art. 105, RED interconnection objective, FuelEU RFNBO multiplier).
- **Quantitative means any number that measures the outcome** — monetary amounts (EUR 65bn SCF), EUR/tonne penalties, and month-count deadlines all count. But a chapeau quote whose numbers sit in unquoted sub-points must not claim to be quantitative — re-extract the sub-points instead.
- **Cite the exact provision.** “Article N” is sharpened to the paragraph/point; text an act inserts into *other* legislation is labelled as an amendment (the Climate Law Art. 13 insertions into Reg. 2018/1999 previously produced citations of non-existent articles); communication rows cite numbered sections, not “Body”.

4.2 · Climate relevance is decided by mechanism, not keyword

Keyword tallies produced both misses (all 13 RefuelEU SAF mandates, AFIR EV-charging rows, CAP eco-schemes tagged *none*) and false positives (chemical “emissions” in the Water Framework Directive, industrial “resilience” in NZIA). The register now decides by the **mechanism the target works through**, and writes the argument — quoting the words that fired it — into its own column so a reader can check it:

Mitigation mechanisms	Adaptation mechanisms
M1 cuts GHGs directly	A1 reduces exposure/vulnerability of people, assets, infrastructure
M2 increases removals or protects the sink	A2 keeps ecosystems and soils able to buffer climate impacts
M3 cuts energy demand	A3 raises adaptive capacity (risk assessment, early warning, planning, insurance)
M4 switches supply or fuels to low-carbon sources	A4 protects health from climate-sensitive hazards
M5 builds enabling conditions (infrastructure, permitting, finance, standards)	A5 seizes an opportunity created by a changing climate
M6 cuts material and product emissions through circularity	

Reviewer-taught calls now encoded: air-pollution and transport-noise targets are **adaptation via A4** (public health); climate-expenditure shares tracking both intervention types are **both**; tree-planting carries the land-sink

mitigation argument; RED building-code rows are mitigation-only. Where a human review set a call explicitly, that call always wins over the classifier.

4.3 · Sector tagging and its traps

Eight sector/system flags (Energy · Buildings · Agri-food · Transport · Industry · Land & marine ecosystems · Water · Health) are assigned from a per-sector vocabulary matched on word boundaries, with explicit **veto phrases** for the traps the training exposed: “soil health” is not Health, “capacity building” is not Buildings, “battery electric vehicle” is not Industry, and F-gas “leak” checks are not Water (five rows had been so tagged). A row whose own words carry no sector vocabulary may inherit its act’s single system, with the inheritance stated in the evidence column. Cross-cutting rows legitimately carry no sector and stay in the register.

4.4 · First-order vs second-order targets

Introduced by the v3 reviewers: **first order** = the overall change the act exists to achieve (subject-matter articles, headline quantified targets and their staged deadlines); **second order** = dependent, complementary or niche targets. 118 human labels are locked as *source: human*; the remaining rows were assigned by calibrated agents with a one-sentence rationale each, so human and machine judgements never silently overwrite each other.

5 · The relevance lens instead of deletion

Expanding to the resilience/health/cohesion acquis pulled in much material peripheral to a climate board. The answer is a boolean **relevant** flag, not removal: a row is relevant if its climate call is not *none*, or it is a quantified, time-bound target — with explicit review overrides winning over the default. The UI opens on the relevant lens; every row remains reachable. Similarly, the 382 rows flagged “Revise target” under NT-1–NT-15 in the v3 pass are **flagged, not removed** — they are the candidate list for the next review round, with the ten human-reviewed acts exempt because their surviving rows are implicitly endorsed.

6 · Source texts must be the current consolidated version

A recurring reviewer theme: figures drifted because the corpus held pre-consolidation texts. The v3 pass swept all 49 CELEX-numbered acts against the Publications Office Cellar and replaced 17 outdated texts, re-extracting their targets. Substantive changes this caught:

- **EU Climate Law** (consolidated 2026-04-07): now includes the adopted binding **90 % net-reduction-by-2040** target.
- **RED III** text: 42.5 %/45 % renewables 2030, buildings 49 %, industry +1.6 pp & RFNBO 42 %/60 %, transport 29 % / –14.5 % GHG intensity.
- **CSDDD** (Omnibus I): the climate-transition Article 22 is **repealed** — the act now contributes no targets, and that fact is documented rather than the old rows kept.
- **EU ETS** (LRF 4.3 %/4.4 %, maritime, ETS2), CO₂ cars (2035 100 % zero-emission), Effort Sharing (–40 %), Waste Framework (55/60/65 % recycling), and ten others.

Lessons: communications are never consolidated (13 acts — skip them); an act named by a reviewer (F-gas) is checked by name even if believed current; acts whose currency cannot be conclusively determined are noted in the sweep audit rather than silently trusted; and reviewer decisions are carried onto re-extracted rows whose content hash changed, so a text refresh never discards human judgement.

7 · Process lessons: how corrections stay trustworthy

- **Generalise every reviewer judgement.** The reviewers said their mark-up was not exhaustive; each flagged row became a written rule applied to all rows. Keyword-flagging reached only 190 of 1 018 rows — the

second pass read every row against its EUR-Lex source (ten parallel agents), and found whole error classes the keywords missed.

- **Propose → adversarially verify.** No agent finding is applied on its own say-so. Verifiers materially changed outcomes — restoring quantitative typing wrongly downgraded and overturning an unjustified drop (TEN-E Art. 1(1)).
- **Re-check claims mechanically where possible.** DUP claims by text containment; NT-1–4 claims only stand where the source actually shows the disqualifying chapeau; every proposed quote re-validated as a verbatim substring, then validated again by the build.
- **Run a completeness check after every apply.** One application silently lost four rows (re-extractions over the then-900-character ceiling were rejected while the rows they superseded were already dropped). “Every candidate must appear in the rebuilt dataset” is now a standing check.
- **Every removal is reversible and named.** All corrections live in *scripts/policy-targets-overrides.json* with a per-entry rule, reason and pass tag, mirrored in the workbook's *Removed rows* sheet — any call can be reversed by deleting one entry and rebuilding (*npm run build:policy-targets*).
- **Row ids are stable content hashes,** so reviewers' workbook mark-ups stay attached to the right rows across regenerations — the prerequisite for this whole workflow (645/646 v3 Revision rows matched their dataset ids).
- **Human labels are locked.** Anything a reviewer set carries *source: human* and is never overwritten by a classifier; agent assignments carry a rationale.

8 · Known limits — where the next reviewer should look

- The NT rules generalise a finite set of reviewer judgements. Coverage is no longer the weak point (every row has been read against its source); the **rules themselves** are — the plan-content (NT-1) and criteria (NT-3) borderlines are where a second reviewer is most likely to disagree.
- Rows deliberately kept but queried: the Renovation Wave and Fit-for-55 restatements of the RRF/MFF climate shares, the CRMA Art. 11 bare duration points, the Water Resilience cohesion-financing package.
- Sector and climate classification is lexical — it reads the target text and its provision heading, not the whole act, so a target whose subject is only implied by its parent provision can be under-tagged.
- Duplicate detection is token overlap — it finds restatements, not two targets that mean the same thing in entirely different words.
- The 382 “Revise target” flags are the ready-made agenda for the next review round.

Current register state (post v3, August 2026): **819 targets across 61 acts, 544 in the default (transition) lens.** Rebuild with *npm run build:policy-targets*; regenerate the reviewer workbook with *python3 scripts/export-policy-targets-workbook.py*.